

# Getting started with PostgreSQL

Gavin Sherry

`gavin@alcove.com.au`

Alcove Systems Engineering

January 16, 2007

# Outline

1 **SELECT 'Hello World'**

2 Hiccups

3 Lets do something

4 Some boring stuff

5 Useful stuff

# Installing PostgreSQL and start

- Just do it through your package manager!

## Fedora/RHEL/CentOS

```
# yum install postgresql
# service postgresql start
```

## Debian/etc

```
# apt-get search postgresql
# /etc/init.d/postgresql-<version> start
```

## Windows installer

- Just run the .msi
- <http://www.postgresql.org/download/>

## Source

```
# ./configure ; make ; make install
# initdb /path/to/data
# pg_ctl -D /path/to/data start
```

It's that simple!

# So, lets go

## Example

```
$ psql postgres
Welcome to psql 8.2.1, the PostgreSQL interactive terminal.

Type:  \copyright for distribution terms
       \h for help with SQL commands
       \? for help with psql commands
       \g or terminate with semicolon to execute query
       \q to quit

postgres=# \d
No relations found.
postgres=# create table foo (i int, t text);
CREATE TABLE
postgres=# insert into foo values(1, 'Hello World');
INSERT 0 1
postgres=# select * from foo;
 i |      t
---+-----
  1 | Hello World
(1 row)
```

# Outline

1 SELECT 'Hello World'

2 **Hiccups**

3 Lets do something

4 Some boring stuff

5 Useful stuff

# FATAL...?

## What's this?

psql: FATAL: IDENT authentication failed for user "gavin"

- Some packages tighten up default security
- What to do:  

```
su - postgres  
createuser gavin
```
- Alternatively, you could relax the authentication – but I won't tell you how!

# Exploring

## Example

```
# select version();
               version
-----
PostgreSQL 8.2.1 ...

# select table_catalog, table_name, table_type from
       information_schema.tables where
       table_schema = 'public';
table_catalog | table_name | table_type
-----+-----+-----
postgres     | foo        | BASE TABLE
(3 rows)
```

# More exploring

## Example

```
# select column_name, data_type from
    information_schema.columns
    where table_schema='public' and
    table_name = 'foo';
```

column_name	data_type
i	integer
t	text

(2 rows)



# Outline

1 SELECT 'Hello World'

2 Hiccups

**3 Lets do something**

4 Some boring stuff

5 Useful stuff

# A little project

- Say you want to develop a blog – “everyone’s doing it”
- You want to:
  - 1 Store articles
  - 2 Relate articles to blog users
  - 3 Keep statistics on blog posts

# A little project

- Say you want to develop a blog – “everyone’s doing it”
- You want to:
  - 1 Store articles
  - 2 Relate articles to blog users
  - 3 Keep statistics on blog posts

# A little project

- Say you want to develop a blog – “everyone’s doing it”
- You want to:
  - 1 Store articles
  - 2 Relate articles to blog users
  - 3 Keep statistics on blog posts

# A little project

- Say you want to develop a blog – “everyone’s doing it”
- You want to:
  - 1 Store articles
  - 2 Relate articles to blog users
  - 3 Keep statistics on blog posts

# blogdb – version 1

## SQL

```
-- Create a table of users (this is a comment, btw)
CREATE TABLE users (usrid int, username text, email text);

-- insert a user (SQL isn't case sensitive)
insert into users values(1, 'Gavin',
                        'gavin@alcove.com.au');

-- create the articles table
CREATE TABLE articles (artid int, title text, body text,
                        dt timestamp, usrid int);

-- add a post
insert into articles values(1, 'First post', 'Yay',
                           current_timestamp, 1);

-- get article count
SELECT count(*) FROM articles;
```

# blogdb – version 1

## SQL

```
-- Create a table of users (this is a comment, btw)
CREATE TABLE users (usrid int, username text, email text);

-- insert a user (SQL isn't case sensitive)
insert into users values(1, 'Gavin',
    'gavin@alcove.com.au');

-- create the articles table
CREATE TABLE articles (artid int, title text, body text,
    dt timestamp, usrid int);

-- add a post
insert into articles values(1, 'First post', 'Yay',
    current_timestamp, 1);

-- get article count
SELECT count(*) FROM articles;
```

# blogdb – version 1

## SQL

```
-- Create a table of users (this is a comment, btw)
CREATE TABLE users (usrid int, username text, email text);

-- insert a user (SQL isn't case sensitive)
insert into users values(1, 'Gavin',
                        'gavin@alcove.com.au');

-- create the articles table
CREATE TABLE articles (artid int, title text, body text,
                        dt timestamp, usrid int);

-- add a post
insert into articles values(1, 'First post', 'Yay',
                           current_timestamp, 1);

-- get article count
SELECT count(*) FROM articles;
```



# blogdb – version 1

## SQL

```
-- Create a table of users (this is a comment, btw)
CREATE TABLE users (usrid int, username text, email text);

-- insert a user (SQL isn't case sensitive)
insert into users values(1, 'Gavin',
    'gavin@alcove.com.au');

-- create the articles table
CREATE TABLE articles (artid int, title text, body text,
    dt timestamp, usrid int);

-- add a post
insert into articles values(1, 'First post', 'Yay',
    current_timestamp, 1);

-- get article count
SELECT count(*) FROM articles;
```

# blogdb – version 1

## SQL

```
-- Create a table of users (this is a comment, btw)
CREATE TABLE users (usrid int, username text, email text);

-- insert a user (SQL isn't case sensitive)
insert into users values(1, 'Gavin',
    'gavin@alcove.com.au');

-- create the articles table
CREATE TABLE articles (artid int, title text, body text,
    dt timestamp, usrid int);

-- add a post
insert into articles values(1, 'First post', 'Yay',
    current_timestamp, 1);

-- get article count
SELECT count(*) FROM articles;
```

# What's wrong with version 1?

- No data integrity
- No permissions
- `users` is a common table name – what about other applications with a `users` table?

# What's wrong with version 1?

- No data integrity
- No permissions
- `users` is a common table name – what about other applications with a `users` table?

# What's wrong with version 1?

- No data integrity
- No permissions
- `users` is a common table name – what about other applications with a `users` table?

# What's wrong with version 1?

- No data integrity
- No permissions
- `users` is a common table name – what about other applications with a `users` table?

# What's wrong with version 1?

- No data integrity
- No permissions
- `users` is a common table name – what about other applications with a `users` table?

# Data integrity

- There are levels of integrity
- Referential integrity
  - Enforce a relationship
- Guard against duplication
  - Each row must have a key
- Some data must look a certain way – like dates and times
- Some data cannot be 'undefined' or `NULL`



# blogdb – version 2

## SQL

```
CREATE TABLE users (usrid serial primary key,  
    username text NOT NULL,  
    email text NOT NULL,  
    unique(username));
```

```
INSERT INTO users (username, email) values(1, 'Gavin',  
    'gavin@alcove.com.au');
```

```
CREATE TABLE articles (artid serial primary key,  
    title text NOT NULL,  
    body text NOT NULL,  
    dt timestamp default current_timestamp,  
    usrid int references users(usrid));
```

# blogdb – version 2

## SQL

```
CREATE TABLE users (usrid serial primary key,  
    username text NOT NULL,  
    email text NOT NULL,  
    unique(username));
```

```
INSERT INTO users (username, email) values(1, 'Gavin',  
    'gavin@alcove.com.au');
```

```
CREATE TABLE articles (artid serial primary key,  
    title text NOT NULL,  
    body text NOT NULL,  
    dt timestamp default current_timestamp,  
    usrid int references users(usrid));
```

# blogdb – version 2

## SQL

```
CREATE TABLE users (usrid serial primary key,  
    username text NOT NULL,  
    email text NOT NULL,  
    unique(username));
```

```
INSERT INTO users (username, email) values(1, 'Gavin',  
    'gavin@alcove.com.au');
```

```
CREATE TABLE articles (artid serial primary key,  
    title text NOT NULL,  
    body text NOT NULL,  
    dt timestamp default current_timestamp,  
    usrid int references users(usrid));
```

# Permissions and privileges

- We speak of CREATE ROLE and GRANT

## SQL

```
CREATE ROLE blogread LOGIN;  
CREATE ROLE blogwrite LOGIN;  
...  
GRANT SELECT ON articles,users  
    TO blogread;  
GRANT ALL ON articles,users  
    TO blobwrite;
```



# Name space issues

- This is why SQL supports “schemas”

## SQL

```
CREATE SCHEMA blog;  
SET search_path = blog;  
CREATE ...  
  
SELECT count(*) FROM blog.articles;
```

## Further exercises

- Data integrity for email addresses – See `CREATE DOMAIN`
- Speed up `count ( * )` in the presence of a large `articles` table – see the advanced lecture today

# Outline

1 SELECT 'Hello World'

2 Hiccups

3 Lets do something

**4 Some boring stuff**

5 Useful stuff

# More on PostgreSQL

- What's with the name? → **Post** (in) **gres**, with **SQL** support
  - Hey, it used to be named POSTQUEL
  - Then again, it used to have a hell of a lot of Lisp in it
- Academic project in the 80s, one of **Michael Stonebraker's** projects at Berkeley
- BSD license – do what you want with the source (don't sue us!)
- Now developed by a group of enthusiasts





# Buzzword compliance

## What have we got?

- JDBC, .NET, ODBC, PHP, Python, Ruby, ...
- Replication
- Views, Indexes, Schemas, Foreign keys
- Partitioning
- Two phase commit
- Transactions and nested transactions!
- Functions/stored procedures in SQL, Java, Ruby, Python, and much more
- Online backups
- ...

# Buzzword compliance

## What have we got?

- JDBC, .NET, ODBC, PHP, Python, Ruby, ...
- Replication
- Views, Indexes, Schemas, Foreign keys
- Partitioning
- Two phase commit
- Transactions and nested transactions!
- Functions/stored procedures in SQL, Java, Ruby, Python, and much more
- Online backups
- ...

# Buzzword compliance

## What have we got?

- JDBC, .NET, ODBC, PHP, Python, Ruby, ...
- Replication
- Views, Indexes, Schemas, Foreign keys
- Partitioning
- Two phase commit
- Transactions and nested transactions!
- Functions/stored procedures in SQL, Java, Ruby, Python, and much more
- Online backups
- ...

# Buzzword compliance

## What have we got?

- JDBC, .NET, ODBC, PHP, Python, Ruby, ...
- Replication
- Views, Indexes, Schemas, Foreign keys
- Partitioning
- Two phase commit
- Transactions and nested transactions!
- Functions/stored procedures in SQL, Java, Ruby, Python, and much more
- Online backups
- ...

# Buzzword compliance

## What have we got?

- JDBC, .NET, ODBC, PHP, Python, Ruby, ...
- Replication
- Views, Indexes, Schemas, Foreign keys
- Partitioning
- Two phase commit
- Transactions and nested transactions!
- Functions/stored procedures in SQL, Java, Ruby, Python, and much more
- Online backups
- ...

# Buzzword compliance

## What have we got?

- JDBC, .NET, ODBC, PHP, Python, Ruby, ...
- Replication
- Views, Indexes, Schemas, Foreign keys
- Partitioning
- Two phase commit
- Transactions and nested transactions!
- Functions/stored procedures in SQL, Java, Ruby, Python, and much more
- Online backups
- ...

# Buzzword compliance

## What have we got?

- JDBC, .NET, ODBC, PHP, Python, Ruby, ...
- Replication
- Views, Indexes, Schemas, Foreign keys
- Partitioning
- Two phase commit
- Transactions and nested transactions!
- Functions/stored procedures in SQL, Java, Ruby, Python, and much more
- Online backups
- ...

# Buzzword compliance

## What have we got?

- JDBC, .NET, ODBC, PHP, Python, Ruby, ...
- Replication
- Views, Indexes, Schemas, Foreign keys
- Partitioning
- Two phase commit
- Transactions and nested transactions!
- Functions/stored procedures in SQL, Java, Ruby, Python, and much more
- Online backups
- ...



# Buzzword compliance

## What have we got?

- JDBC, .NET, ODBC, PHP, Python, Ruby, ...
- Replication
- Views, Indexes, Schemas, Foreign keys
- Partitioning
- Two phase commit
- Transactions and nested transactions!
- Functions/stored procedures in SQL, Java, Ruby, Python, and much more
- Online backups
- ...

# Buzzword compliance

## What have we got?

- JDBC, .NET, ODBC, PHP, Python, Ruby, ...
- Replication
- Views, Indexes, Schemas, Foreign keys
- Partitioning
- Two phase commit
- Transactions and nested transactions!
- Functions/stored procedures in SQL, Java, Ruby, Python, and much more
- Online backups
- ...

# Buzzword compliance

## What have we got?

- JDBC, .NET, ODBC, PHP, Python, Ruby, ...
- Replication
- Views, Indexes, Schemas, Foreign keys
- Partitioning
- Two phase commit
- Transactions and nested transactions!
- Functions/stored procedures in SQL, Java, Ruby, Python, and much more
- Online backups
- ...

# Buzzword compliance

## What have we got?

- JDBC, .NET, ODBC, PHP, Python, Ruby, ...
- Replication
- Views, Indexes, Schemas, Foreign keys
- Partitioning
- Two phase commit
- Transactions and nested transactions!
- Functions/stored procedures in SQL, Java, Ruby, Python, and much more
- Online backups
- ...

# Buzzword compliance

## What have we got?

- JDBC, .NET, ODBC, PHP, Python, Ruby, ...
- Replication
- Views, Indexes, Schemas, Foreign keys
- Partitioning
- Two phase commit
- Transactions and nested transactions!
- Functions/stored procedures in SQL, Java, Ruby, Python, and much more
- Online backups
- ...

# Buzzword compliance

## What have we got?

- JDBC, .NET, ODBC, PHP, Python, Ruby, ...
- Replication
- Views, Indexes, Schemas, Foreign keys
- Partitioning
- Two phase commit
- Transactions and nested transactions!
- Functions/stored procedures in SQL, Java, Ruby, Python, and much more
- Online backups
- ...

# Buzzword compliance

## What have we got?

- JDBC, .NET, ODBC, PHP, Python, Ruby, ...
- Replication
- Views, Indexes, Schemas, Foreign keys
- Partitioning
- Two phase commit
- Transactions and nested transactions!
- Functions/stored procedures in SQL, Java, Ruby, Python, and much more
- Online backups
- ...

# Buzzword compliance

## What have we got?

- JDBC, .NET, ODBC, PHP, Python, Ruby, ...
- Replication
- Views, Indexes, Schemas, Foreign keys
- Partitioning
- Two phase commit
- Transactions and nested transactions!
- Functions/stored procedures in SQL, Java, Ruby, Python, and much more
- Online backups
- ...



# Buzzword compliance

## What have we got?

- JDBC, .NET, ODBC, PHP, Python, Ruby, ...
- Replication
- Views, Indexes, Schemas, Foreign keys
- Partitioning
- Two phase commit
- Transactions and nested transactions!
- Functions/stored procedures in SQL, Java, Ruby, Python, and much more
- Online backups
- ...

# Buzzword compliance

## What have we got?

- JDBC, .NET, ODBC, PHP, Python, Ruby, ...
- Replication
- Views, Indexes, Schemas, Foreign keys
- Partitioning
- Two phase commit
- Transactions and nested transactions!
- Functions/stored procedures in SQL, Java, Ruby, Python, and much more
- Online backups

# Buzzword compliance

## What have we got?

- JDBC, .NET, ODBC, PHP, Python, Ruby, ...
- Replication
- Views, Indexes, Schemas, Foreign keys
- Partitioning
- Two phase commit
- Transactions and nested transactions!
- Functions/stored procedures in SQL, Java, Ruby, Python, and much more
- Online backups
- ...

# Who uses PostgreSQL?

- Businesses, big and small
- Hobbyists and developers
- Big operators like
  - Fujitsu
  - Sun
  - HP
  - Apple
  - Cisco
  - AC Neilson
  - Skype
  - Just about every Australian government department
  - Safeway (North America)
  - .org, .info, .in, ...
  - Sony
- But it's even more popular with small businesses

# What else can you do?

- Report writing: Jasper, Crystal, ...
- Wiki, CMS backend
- All the other OSS jazz – forums, photo galleries, ...
- Bugzilla backend
- LDAP backend
- J2EE backend
- Ruby on rails, out of the box
- GIS – 100% standards compliant!
- Statistical analysis – full integration of 'R'
- Data warehousing

# Outline

1 SELECT 'Hello World'

2 Hiccups

3 Lets do something

4 Some boring stuff

**5 Useful stuff**

# Resources for learning more

- The PostgreSQL manual -  
<http://www.postgresql.org/docs>
- Books (some free) -  
<http://www.postgresql.org/docs/books/>
- GUI console - <http://www.pgadmin.org>
- Mailing lists - <http://www.postgresql.org/>
- IRC - [irc.freenode.net](http://irc.freenode.net), #postgresql
  - My nick is swm
- SydpUG - <http://pugs.postgresql.org/sydpug/>

# Stuff for the boss

- Case studies:

<http://www.postgresql.org/about/casestudies/>

- Featured users:

<http://www.postgresql.org/about/users/>

- A glossy brochure:

<http://alcove.com.au/postgresql.pdf>